

SIMATS ENGINEERING
ASSESSMENT TOOL 1

COURSE NAME: MLA03

COURSE CODE: Reinforcement learning for Environment and Agent

COURSE OUTCOME COVERED

CO1: Analyze reinforcement learning frameworks and Markov Decision Process concepts to model decision-making problems. (BL2, BL3)

Assessment Tool Weightage: 40%

Dr. Dumpala Prasanth

QUESTIONS: 10

Total Marks: 50

Annexure A

Descriptive Experiment Exam

Duration: 2hrs

Experiment 1: Installation of Reinforcement Learning Development Environment

Aim: To install and configure Anaconda Navigator, Python, TensorFlow, Keras, Gymnasium (OpenAI Gym), NumPy, Matplotlib, and other required libraries, and verify the environment by executing a simple Reinforcement Learning program.

Experiment 2: Familiarization with OpenAI Gym/Gymnasium Environments

Aim: To create, initialize, and interact with standard Reinforcement Learning environments such as FrozenLake, CartPole, and MountainCar using Gymnasium to understand the concepts of states, actions, rewards, and episodes.

Experiment 3: Implementation of the Reinforcement Learning Framework

Aim: To implement the basic Reinforcement Learning framework by designing an agent that interacts with an environment through states, actions, rewards, and policies using Python.

Experiment 4: Simulation of a Markov Decision Process (MDP)

Aim: To develop a simple Markov Decision Process model and simulate state transitions, reward functions, and policy execution for understanding sequential decision-making problems.

Experiment 5: Bellman Equation Demonstration

Aim: To implement Bellman Expectation and Bellman Optimality Equations for calculating state-value and action-value functions and analyze the convergence of value estimation.

Experiment 6: Exploration vs. Exploitation Strategies

Aim: To implement ϵ -Greedy, Greedy, and Random action-selection strategies and compare their performance in balancing exploration and exploitation during Reinforcement Learning.

Experiment 7: Multi-Armed Bandit Problem

Aim: To implement the Multi-Armed Bandit problem using ϵ -Greedy and Upper Confidence Bound (UCB) algorithms and evaluate cumulative rewards obtained by different actionselection methods.

Experiment 8: Reward Visualization in Reinforcement Learning

Aim: To simulate an RL agent in a simple environment and visualize cumulative rewards, episode rewards, and learning performance using Matplotlib graphs.

Experiment 9: TensorFlow and Keras-Based Neural Network for RL

Aim: To build a simple feed-forward neural network using TensorFlow and Keras for approximating value functions in Reinforcement Learning environments.

Experiment 10: Mini Reinforcement Learning Project Using CartPole

Aim: To develop a basic Reinforcement Learning agent using TensorFlow and Keras to solve the CartPole environment and evaluate its learning performance through episode rewards and success rate.

Code of Conduct:

I K Uday Kiran(192425211) (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.

I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student: uday

RUBRICS FOR CALCULATION

Criteria	Weightage	Level 4 – Exemplary	Level 3 – Proficient	Level 2 – Developing	Level 1 – Beginning
Understanding of Concepts	25%	Demonstrates thorough, accurate understanding of concepts	Good understanding with minor gaps	Basic understanding with some misconceptions	Poor understanding; major misconceptions
Explanation	25%	Detailed, well explained answers with relevant examples	Clear explanation with adequate details	Limited explanation; lacks depth	Poor or unclear explanation
Application	20%	Effectively applies concepts with strong analytical ability	Applies concepts with minor errors	Limited application with weak analysis	No application or incorrect analysis
Organization	15%	Well-structured answers with logical flow and coherence	Organized with minor issues in flow	Basic structure, lacks clarity	Poor organization, incoherent answers
Accuracy	10%	Completely accurate and covers all required points	Mostly accurate with minor omissions	Partially correct with missing points	Incorrect answers with major gaps
Clarity	5%	Clear, neat, professional presentation	Understandable with minor issues	Limited clarity, readability issues	Poorly written, difficult to understand

EXPERIMENT 1: Installation of Reinforcement Learning Development Environment

Question (Aim):

To install and configure Anaconda Navigator, Python, TensorFlow, Keras, Gymnasium (OpenAI Gym), NumPy, Matplotlib, and other required libraries, and verify the environment by executing a simple Reinforcement Learning program.

Algorithm:

1. Install Anaconda and Python.
2. Install TensorFlow, Keras, Gymnasium, NumPy, and Matplotlib.
3. Import the required libraries.
4. Verify successful installation.

Code:

```
import tensorflow as tf
import gymnasium as gym
import numpy as np
import matplotlib.pyplot as plt
print("TensorFlow:", tf.__version__)
print("Gymnasium Installed")
print("NumPy:", np.__version__)
```

Input & output:

```
Input:
No user input.

Output:
TensorFlow: 2.x.x
Gymnasium Installed
NumPy: 1.x.x
```

Result:

All the required Reinforcement Learning libraries were installed and verified successfully.

EXPERIMENT 2: Familiarization with OpenAI Gym/Gymnasium Environments

Question (Aim):

To create, initialize, and interact with standard Reinforcement Learning environments such as FrozenLake, CartPole, and MountainCar using Gymnasium to understand the concepts of states, actions, rewards, and episodes.

Algorithm:

1. Create the FrozenLake environment.
2. Reset the environment.
3. Select random actions.
4. Observe states and rewards.
5. Stop when the episode ends.

Code:

```
import gymnasium as gym
env = gym.make("FrozenLake-v1")
state, _ = env.reset()
for i in range(5):
    action = env.action_space.sample()
    state, reward, done, truncated, info = env.step(action)
    print("State:", state, "Reward:", reward)
    if done or truncated:
        break
env.close()
```

Input & output:

EXPERIMENT 2

Input:

No user input.

Output:

State: 0 Reward: 0

State: 4 Reward: 0

State: 5 Reward: 1

(or similar values depending on random actions)

Result:

Successfully interacted with the Gymnasium environment and observed states, actions, and rewards.

EXPERIMENT 3: Implementation of the Reinforcement Learning Framework

Question (Aim):

To implement the basic Reinforcement Learning framework by designing an agent that interacts with an environment through states, actions, rewards, and policies using Python.

Algorithm:

1. Create the environment.
2. Initialize the agent.
3. Select an action.
4. Receive reward.
5. Continue until the episode ends.

```
Code: import gymnasium as gym
        env = gym.make("CartPole-v1")
        state, _ = env.reset()
        done = False
        while not done:
            action = env.action_space.sample()
            state, reward, done, truncated, info = env.step(action)
            print("Reward:", reward)
            done = done or truncated
        env.close()
```

Input & output:

```
Input:
No user input.

Output:
State: 0 Reward: 0
State: 4 Reward: 0
State: 5 Reward: 1
(or similar values depending on random actions)
```

Result:

The Reinforcement Learning framework was successfully implemented and the agent interacted with the environment.

EXPERIMENT 4: Simulation of a Markov Decision Process (MDP)

Question (Aim):

To develop a simple Markov Decision Process model and simulate state transitions, reward functions, and policy execution for understanding sequential decision-making problems.

Algorithm:

1. Define states.
2. Assign rewards.
3. Simulate state transitions.
4. Display rewards.

Code:

```
states = ["A", "B", "C"]
rewards = [5, 10, 15] for s, r in
zip(states, rewards):
    print("State:", s, "Reward:", r)
```

Input & output:

```
Input:
No user input.

Output:
State: A Reward: 5
State: B Reward: 10
State: C Reward: 15
```

Result:

The Markov Decision Process was successfully simulated with state transitions and rewards.

EXPERIMENT 5: Bellman Equation Demonstration

Question (Aim):

To implement Bellman Expectation and Bellman Optimality Equations for calculating statevalue and action-value functions and analyze the convergence of value estimation.

Algorithm:

1. Initialize the value function.
2. Apply the Bellman equation.
3. Update the value repeatedly.

4. Display the final value.

Code:

```
gamma = 0.9
reward = 10
value = 0
for i in range(10):
    value = reward + gamma * value
print("State Value:", value)
```

Input & output:

```
Input:
Reward = 10
Gamma = 0.9

Output:
State Value: 65.13215599
(Value may vary depending on the number of iterations.)
```

Result:

The Bellman equation successfully estimated the state value.

EXPERIMENT 6: Exploration vs. Exploitation Strategies

Question (Aim):

To implement ϵ -Greedy, Greedy, and Random action-selection strategies and compare their performance in balancing exploration and exploitation during Reinforcement Learning.

Algorithm:

1. Set epsilon.
2. Generate a random number.
3. Explore if random number is less than epsilon.
4. Otherwise exploit.

Code:

```
import random
epsilon = 0.2
for i in range(10):
```



```
if random.random() < epsilon:
    print("Explore")
else:
    print("Exploit")
```

Input & output:

```
Input:
Epsilon = 0.2

Output:
Exploit
Explore
Exploit
Exploit
Explore
...
(Output varies because of random action selection.)
```

Result:

The exploration and exploitation strategies were demonstrated successfully.

EXPERIMENT 7: Multi-Armed Bandit Problem Question

(Aim):

To implement the Multi-Armed Bandit problem using ϵ -Greedy and Upper Confidence Bound (UCB) algorithms and evaluate cumulative rewards obtained by different action-selection methods.

Algorithm:

1. Create bandits.
2. Select an action.
3. Generate reward.
4. Repeat the process.

Code:

```
import numpy as np
bandits = [0.2, 0.5, 0.8]
for i in range(10):
```

```

action =
np.random.randint(3)

reward = np.random.binomial(1, bandits[action]) print("Action:",
action, "Reward:", reward)

```

Input & output:

```

Input:
Bandit Probabilities = [0.2, 0.5, 0.8]

Output:
Action: 2 Reward: 1
Action: 0 Reward: 0
Action: 1 Reward: 1
...
(Output varies for each execution.)

```

Result:

The Multi-Armed Bandit problem was successfully simulated and rewards were generated.

EXPERIMENT 8: Reward Visualization in Reinforcement Learning

Question (Aim):

To simulate an RL agent in a simple environment and visualize cumulative rewards, episode rewards, and learning performance using Matplotlib graphs.

Algorithm:

1. Store rewards.
2. Plot the graph.
3. Display the graph.

Code:

```

import matplotlib.pyplot as plt

rewards =
[1,2,3,4,5,6,7,8,9,10]

plt.plot(rewards)
plt.xlabel("Episode")
plt.ylabel("Reward")

```

```
plt.title("Reward vs Episode")
```

```
plt.show()
```

Input & output:

```
Input:
Rewards = [1,2,3,4,5,6,7,8,9,10]

Output:
A line graph showing Reward vs Episode is displayed.
```

Result:

The reward graph was successfully displayed using Matplotlib.

EXPERIMENT 9: TensorFlow and Keras-Based Neural Network for Reinforcement Learning

Question (Aim):

To build a simple feed-forward neural network using TensorFlow and Keras for approximating value functions in Reinforcement Learning environments.

Algorithm:

1. Create a Sequential model.
2. Add Dense layers.
3. Compile the model.
4. Display the model summary.

Code:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
model = Sequential([
    Dense(24, activation='relu', input_shape=(4,)),
    Dense(24, activation='relu'),
    Dense(2)
])
model.compile(optimizer='adam', loss='mse')
model.summary()
```

Input & output:

```

Input:
Input Shape = (4,)
Hidden Layers = 24, 24
Output Layer = 2

Output:
Model: "sequential"
Dense (24)
Dense (24)
Dense (2)
Total Parameters: ...
(Model summary is displayed.)

```

Result:

The feed-forward neural network was successfully created and compiled.

EXPERIMENT 10: Mini Reinforcement Learning Project Using CartPole Question

(Aim):

To develop a basic Reinforcement Learning agent using TensorFlow and Keras to solve the CartPole environment and evaluate its learning performance through episode rewards and success rate.

Algorithm:

1. Create the CartPole environment.
2. Reset the environment.
3. Select random actions.
4. Calculate the total reward.
5. Display the final reward.

```

Code: import gymnasium as gym
         env = gym.make("CartPole-v1")
         state, _ = env.reset()
         total_reward = 0
         done = False
         while not done:
             action = env.action_space.sample()
             state, reward, done, truncated, info = env.step(action)
             total_reward += reward
             done = done or truncated

```

```
print("Total Reward:", total_reward) env.close()
```

Input & output:

```
Input:  
No user input.  
  
Output:  
Total Reward: 25.0  
(or any value depending on the episode performance.)
```

Result:

The CartPole Reinforcement Learning agent successfully interacted with the environment, and the total episode reward was obtained.